

Algunos ejemplos en R para Exploración de Datos

Para más detalles, ver *Data Mining and Predictive Analytics – Daniel T. Larose y Chantal D. Larose – Wiley 2015*

Leer los datasets Cars y Cars2 datasets:

Reemplazar "C:/ ... /" con la localización exacta del archivo que se desea abrir

```
cars <- read.csv("C:/ ... /cars.txt", stringsAsFactors = FALSE)
cars2 <- read.csv("C:/ ... /cars2.txt", stringsAsFactors = FALSE)
```

(cars cuenta con 261 observaciones y 8 variables, mientras que cars2 cuenta con 263 observaciones y 8 variables)

Datos Faltantes

Ver cuatro variables del conjunto cars

```
cars.4var <- cars[,c(1,3,4,8)]
head(cars.4var)
```

261 observaciones y 4 variables

	mpg	cubicinches	hp	brand
1	14.0	350	165	US.
2	31.9	89	71	Europe.
3	17.0	302	140	US.
4	15.0	400	150	US.
5	30.5	98	63	US.
6	23.0	350	125	US.

Convertir a Missing algunos valores

```
cars.4var[2,2] <- cars.4var[4,4] <- NA
head(cars.4var)
```

	mpg	cubicinches	hp	brand
1	14.0	350	165	US.
2	31.9	NA	71	Europe.
3	17.0	302	140	US.
4	15.0	400	150	<NA>
5	30.5	98	63	US.
6	23.0	350	125	US.

Reemplazar valores faltantes con constantes

```
cars.4var[2,2] <- 0
cars.4var[4,4] <- "Missing"
head(cars.4var)
```

	mpg	cubicinches	hp	brand
1	14.0	350	165	US.
2	31.9	0	71	Europe.
3	17.0	302	140	US.
4	15.0	400	150	Missing
5	30.5	98	63	US.
6	23.0	350	125	US.

Reemplazar valores con la media y la moda

```
cars.4var[2,2] <- mean(na.omit(cars.4var$cubicinches))
our_table <- table(cars.4var$brand)
our_mode <- names(our_table)[our_table == max(our_table)]
cars.4var[4,4] <- our_mode
head(cars.4var)
```

	mpg	cubicinches	hp	brand
1	14.0	350.0000	165	US.
2	31.9	200.7625	71	Europe.
3	17.0	302.0000	140	US.
4	15.0	400.0000	150	US.
5	30.5	98.0000	63	US.
6	23.0	350.0000	125	US.

Generar observaciones aleatorias

```
obs_brand <- sample(na.omit(cars.4var$brand), 1)
obs_cubicinches <- sample(na.omit(cars.4var$cubicinches), 1)
cars.4var[2,2] <- obs_cubicinches
cars.4var[4,4] <- obs_brand
head(cars.4var)
```

	mpg	cubicinches	hp	brand
1	14.0	350	165	US.
2	31.9	97	71	Europe.
3	17.0	302	140	US.
4	15.0	400	150	US.
5	30.5	98	63	US.
6	23.0	350	125	US.

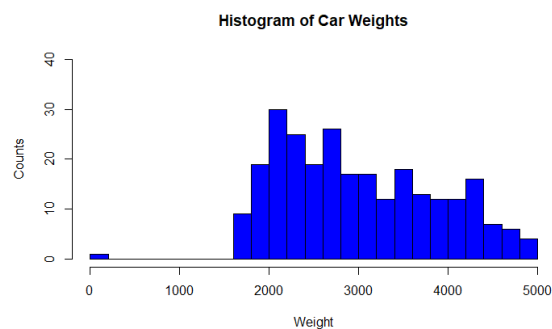
Crear un Histograma

Definir el área de dibujo

```
par(mfrow = c(1,1))
```


```
# Crear el histograma
```

```
hist(cars2$weight, breaks = 30, xlim= c(0, 5000), col = "blue", border = "black", ylim= c(0, 40), xlab = "Weight", ylab = "Counts", main = "Histogram of Car Weights")
```



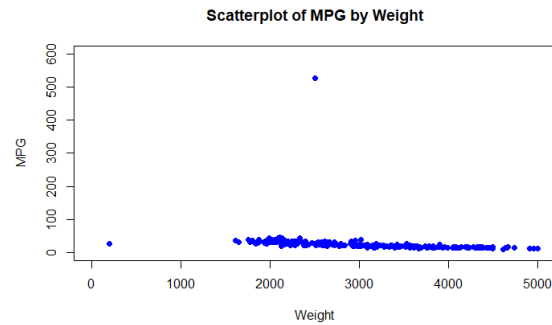
Colocar una línea alrededor del gráfico

```
box(which = "plot", lty = "solid", col = "black")
```



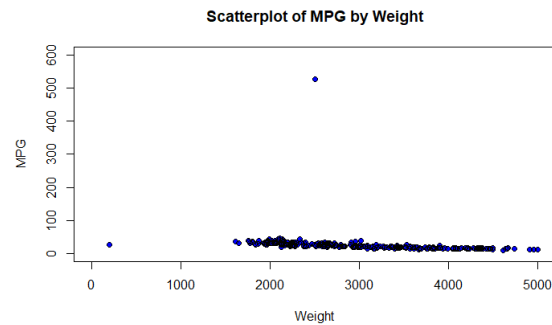
Crear un gráfico de dispersión

```
plot(cars2$weight, cars2$mpg, xlim= c(0, 5000), ylim= c(0, 600), xlab = "Weight", ylab = "MPG", main = "Scatterplot of MPG by Weight", type = "p", pch = 16, col = "blue")
```



#Agregar los puntos como círculos negros sin relleno

```
points(cars2$weight, cars2$mpg, type = "p", col = "black")
```



Estadísticas Descriptivas

```
mean(cars$weight)    # Mean
median(cars$weight)   # Median
length(cars$weight)  # Number of observations
sd(cars$weight)       # Standard deviation
summary(cars$weight)  # Min, Q1, Median, Mean, Q3, Max
```

```
> mean(cars$weight)    # Mean
[1] 3005.49
> median(cars$weight)  # Median
[1] 2835
> length(cars$weight)  # Number of observations
[1] 261
> sd(cars$weight)      # Standard deviation
[1] 852.6456
> summary(cars$weight) # Min, Q1, Median, Mean, Q3, Max
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   1613   2246   2835   3005   3664   4997
```

Transformaciones

Min-max normalization

```
summary(cars$weight)
mi <- min(cars$weight)
ma <- max(cars$weight)
minmax.weight <- (cars$weight - mi)/(ma - mi)
minmax.weight
```

```
[1] 0.76713948 0.09219858 0.54255319 0.63475177 0.12943262 0.67582742 0.81264775 0.79757683
 [9] 0.56648936 0.12913712 0.18676123 0.16991726 0.74704492 0.60608747 0.52245863 0.81176123
[17] 0.12027187 0.53782506 0.19651300 0.24143026 0.11997636 0.29166667 0.41341608 0.72133570
[25] 0.10490544 0.05378251 0.04728132 0.89952719 0.57949173 0.22665485 0.15277778 0.49202128
[33] 0.26388889 0.30585106 0.10992908 0.53250591 0.06560284 0.32712766 0.69208038 0.48817967
[41] 0.67494090 0.89391253 0.60047281 0.56353428 0.43528369 0.30200946 0.63297872 0.23108747
...
```

Z-score standarization

```
m <- mean(cars$weight); s <- sd(cars$weight)
z.weight <- (cars$weight - m)/s
z.weight
length(cars$weight)
```

```
[1] 1.411500375 -1.267221019 0.520157004 0.886076914 -1.119445671 1.049098925 1.592114690
 [8] 1.532300858 0.615155442 -1.120618491 -0.891918547 -0.958769300 1.331748600 0.772313352
[15] 0.440405229 1.588596229 -1.155803098 0.501391880 -0.853215480 -0.674946806 -1.156975918
[22] -0.475567367 0.007634565 1.229713240 -1.216789750 -1.419687649 -1.445489694 1.936923836
[29] 0.666759532 -0.733587817 -1.026792873 0.319604745 -0.585812469 -0.419271997 -1.196851806
[36] 0.480281116 -1.372774840 -0.334828940 1.113604038 0.304358082 1.045580465 1.914640252
...
```

Decimal scaling

```
max(abs(cars$weight)) # 4 digits (el Maximo es 4997)
d.weight <- cars$weight/(10^4); d.weight
```

```
[1] 0.4209 0.1925 0.3449 0.3761 0.2051 0.3900 0.4363 0.4312 0.3530 0.2050 0.2245 0.2188 0.4141
[14] 0.3664 0.3381 0.4360 0.2020 0.3433 0.2278 0.2430 0.2019 0.2600 0.3012 0.4054 0.1968 0.1795
[27] 0.1773 0.4657 0.3574 0.2380 0.2130 0.3278 0.2506 0.2648 0.1985 0.3415 0.1835 0.2720 0.3955
[40] 0.3265 0.3897 0.4638 0.3645 0.3520 0.3086 0.2635 0.3755 0.2395 0.1940 0.3060 0.4464 0.3190
...
```

Histogramas lado a lado

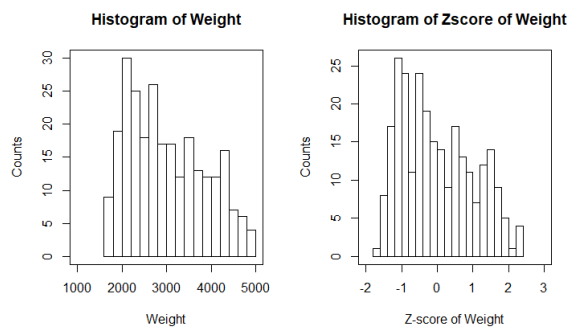
```
par(mfrow = c(1,2))
```

Crear dos histogramas

```
hist(cars$weight, breaks = 20,
xlim = c(1000, 5000),
main = "Histogram of Weight",
xlab = "Weight",
```

```
ylab = "Counts")
box(which = "plot",
lty = "solid",
col = "black")

hist(z.weight,
breaks = 20, xlim = c(-2, 3),
main = "Histogram of Zscore of Weight",
xlab = "Z-score of Weight",
ylab = "Counts")
box(which = "plot",
lty = "solid",
col = "black")
```



Sesgo

```
(3*(mean(cars$weight) - median(cars$weight)))/sd(cars$weight)
(3*(mean(z.weight) - median(z.weight)))/sd(z.weight)
```

(es el mismo valor en ambos casos...)

Transformations for Normality

```
sqrt.weight <- sqrt(cars$weight) # Square root
sqrt.weight_skew <- (3*(mean(sqrt.weight) - median(sqrt.weight)))/sd(sqrt.weight)
ln.weight <- log(cars$weight) # Natural log
ln.weight_skew <- (3*(mean(ln.weight) - median(ln.weight)))/sd(ln.weight)
invsqrt.weight <- 1 / sqrt(cars$weight) # Inverse square root
invsqrt.weight_skew <- (3*(mean(invsqrt.weight) - median(invsqrt.weight)))/sd(invsqrt.weight)
```

Sesgo para transformación raíz cuadrada = 0.4028


Sesgo para Logaritmo Natural = 0.1956

Sesgo para Raíz Cuadrada Inversa = 0.0154

Histograma con curva de distribución normal


```
par(mfrow=c(1,1))
x <- rnorm(1000000, mean = mean(invsqrt.weight), sd = sd(invsqrt.weight))
hist(invsqrt.weight, breaks = 30, xlim = c(0.0125, 0.0275), col = "lightblue", prob =
TRUE, border = "black", xlab="Inverse Square Root of Weight", ylab = "Counts", main =
"Histogram of Inverse Square Root of Weight")
box(which = "plot", lty = "solid", col="black")
```

```
# Densidad Normal
lines(density(x), col="red")
```



Normal Q-Q Plot

```
qqnorm(invsqrt.weight, datax = TRUE, col = "red", ylim = c(0.01, 0.03), main = "Normal
Q-Q Plot of Inverse Square Root of Weight")
qqline(invsqrt.weight, col = "blue", datax = TRUE)
```



Detransformar datos

```
# Transformar x empleando y = 1 / sqrt(x)
x <- cars$weight[1]; y <- 1 / sqrt(x)

# Detransformar x empleando x = 1 / (y)^2
detransformedx <- 1 / y^2
x; y; detransformedx

> x; y; detransformedx
[1] 4209
[1] 0.01541383
[1] 4209
```

Crear variables indicadoras

```
north_flag <- east_flag <- south_flag <- c(rep(NA, 10))
```

```
region <- c(rep(c("north", "south", "east", "west"),2), "north", "south")
# Change the region variable to indicators
for (i in 1:length(region)) {
  if(region[i] == "north") north_flag[i] = 1 else north_flag[i] = 0
  if(region[i] == "east") east_flag[i] = 1 else east_flag[i] = 0
  if(region[i] == "south") south_flag[i] = 1 else south_flag[i] = 0 }
north_flag; east_flag; south_flag

north_flag; east_flag; south_flag
[1] 1 0 0 0 1 0 0 0 1 0
[1] 0 0 1 0 0 0 1 0 0 0
[1] 0 1 0 0 0 1 0 0 0 1
```

Index fields

```
# Data frames have an index field;
# the left-most column cars
cars[order(cars$mpg),]

# For vectors or matrices,
# add a column to act as an index field
x <- c(1,1,3:1,1:4,3);
y <- c(9,9:1)
z <- c(2,1:9)
matrix <- t(rbind(x,y,z)); matrix
indexed_m <- cbind(c(1:length(x)), matrix); indexed_m
indexed_m[order(z),]

> matrix <- t(rbind(x,y,z)); matrix
      x y z
[1,] 1 9 2
[2,] 1 9 1
[3,] 3 8 2
[4,] 2 7 3
[5,] 1 6 4
[6,] 1 5 5
[7,] 2 4 6
[8,] 3 3 7
[9,] 4 2 8
[10,] 3 1 9
> indexed_m <- cbind(c(1:length(x)), matrix); indexed_m
      x y z
[1,] 1 1 9 2
[2,] 2 1 9 1
[3,] 3 3 8 2
[4,] 4 2 7 3
[5,] 5 1 6 4
[6,] 6 1 5 5
[7,] 7 2 4 6
[8,] 8 3 3 7
[9,] 9 4 2 8
[10,] 10 3 1 9
> indexed_m[order(z),]
      x y z
```




```
[1,] 2 1 9 1
[2,] 1 1 9 2
[3,] 3 3 8 2
[4,] 4 2 7 3
[5,] 5 1 6 4
[6,] 6 1 5 5
[7,] 7 2 4 6
[8,] 8 3 3 7
[9,] 9 4 2 8
[10,] 10 3 1 9
```

Duplicate records

```
# For number of duplicate records, use anyDuplicated
anyDuplicated(cars)
# To examine each record, use Duplicated
duplicated(cars)
# 'True': record is a duplicate,
# 'False': record is not a duplicate
# Let's duplicate the first record
new.cars <- rbind(cars, cars[1,])
# Check for duplicates
anyDuplicated(new.cars)
# The 262nd record is a duplicate
duplicated(new.cars)
```

Código en R para la exploración de datos del caso 2

Leer el conjunto de datos

```
churn <- read.csv(file = "C:/ ... /churn.txt", stringsAsFactors=TRUE)
# Mostrar los primeros diez registros
churn[1:10,]
```

	State	Account.Length	Area.Code	Phone	Int.l.Plan	VMail.Plan	VMail.Message	Day.Mins	Day.Calls
1	KS	128	415	382-4657	no	yes	25	265.1	110
2	OH	107	415	371-7191	no	yes	26	161.6	123
3	NJ	137	415	358-1921	no	no	0	243.4	114
4	OH	84	408	375-9999	yes	no	0	299.4	71
5	OK	75	415	330-6626	yes	no	0	166.7	113
6	AL	118	510	391-8027	yes	no	0	223.4	98
7	MA	121	510	355-9993	no	yes	24	218.2	88
8	MO	147	415	329-9001	yes	no	0	157.0	79
9	LA	117	408	335-4719	no	no	0	184.5	97
10	WV	141	415	330-8173	yes	yes	37	258.6	84
	Day.Charge	Eve.Mins	Eve.Calls	Eve.Charge	Night.Mins	Night.Calls	Night.Charge	Intl.Mins	Intl.Calls
1s									
1	45.07	197.4	99	16.78	244.7	91	11.01	10.0	
3									
2	27.47	195.5	103	16.62	254.4	103	11.45	13.7	
3									
3	41.38	121.2	110	10.30	162.6	104	7.32	12.2	
5									
4	50.90	61.9	88	5.26	196.9	89	8.86	6.6	
7									
5	28.34	148.3	122	12.61	186.9	121	8.41	10.1	
3									
6	37.98	220.6	101	18.75	203.9	118	9.18	6.3	
6									
7	37.09	348.5	108	29.62	212.6	118	9.57	7.5	
7									
8	26.69	103.1	94	8.76	211.8	96	9.53	7.1	
6									
9	31.37	351.6	80	29.89	215.8	90	9.71	8.7	
4									
10	43.96	222.0	111	18.87	326.4	97	14.69	11.2	
5									
	Intl.Charge	CustServ.Calls	Churn.						
1	2.70		1 False.						
2	3.70		1 False.						
3	3.29		0 False.						
4	1.78		2 False.						
5	2.73		3 False.						
6	1.70		0 False.						
7	2.03		3 False.						
8	1.92		0 False.						
9	2.35		1 False.						
10	3.02		0 False.						

Sumarizar la variable CHURN

```
sum.churn <- summary(churn$Churn)
sum.churn
```

```
> sum.churn
False.  True.
 2850   483
```

Calcular la proporción de CHURNERS en el total de registros

```
prop.churn <- sum(churn$Churn == "True.") / length(churn$Churn)
prop.churn
```

```
> prop.churn
[1] 0.1449145
```

Grafico de Barras de la variable Churn

```
barplot(sum.churn, ylim = c(0, 3000), main = "Bar Graph of Churners and Non-Churners",
col = "lightblue")
box(which = "plot", lty = "solid", col="black")
```




Tabla doble entrada para churners e International Plan

```
counts <- table(churn$Churn, churn$Int.l.Plan, dnn=c("Churn", "International Plan"))
counts
```

```
> counts
      International Plan
Churn   no  yes
False. 2664 186
True.   346 137
```

Grafico de Barras Superpuesto Proporciones de churners por International Plan

```
barplot(counts, legend = rownames(counts), col = c("blue", "red"), ylim = c(0, 3300),
ylab = "Count", xlab = "International Plan", main = "Comparison Bar Chart: Churn
Proportions by International Plan")
box(which = "plot", lty = "solid", col="black")
```

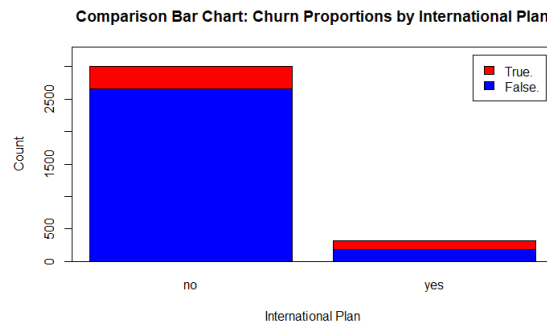


Tabla suma para ambas variables

```
sumtable <- addmargins(counts, FUN = sum)
sumtable
```

```
> sumtable
      International Plan
Churn   no  yes  sum
False. 2664 186 2850
True.   346 137  483
sum     3010 323 3333
```

Crear una tabla de proporciones sobre las filas

```
row.margin <- round(prop.table(counts, margin = 1), 4)*100
row.margin
```

```
> row.margin
      International Plan
Churn   no  yes
False. 93.47  6.53
True.   71.64 28.36
```

Crear una tabla de proporciones sobre las columnas

```
col.margin <- round(prop.table(counts, margin = 2), 4)*100
col.margin
```

```
> col.margin
      International Plan
Churn    no    yes
False.  88.50 57.59
True.   11.50 42.41
```

Grafico de Barras Agrupado, con referencias

```
barplot(counts, col = c("blue", "red"), ylim = c(0, 3300), ylab = "Count", xlab =
"International Plan", main = "Churn Count by International Plan", beside = TRUE)
legend("topright", c(rownames(counts)), col = c("blue", "red"), pch = 15, title =
"Churn")
box(which = "plot", lty = "solid", col="black")
```

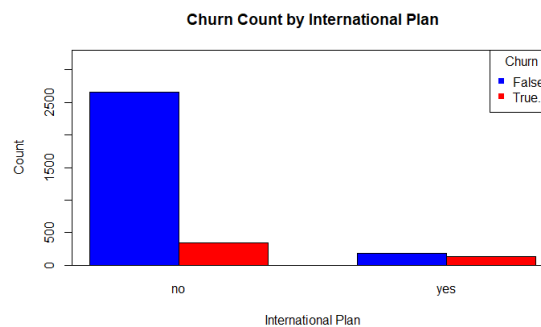
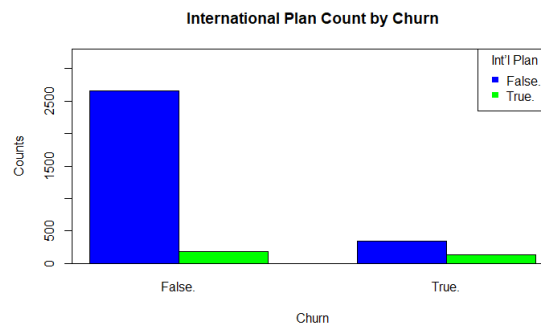


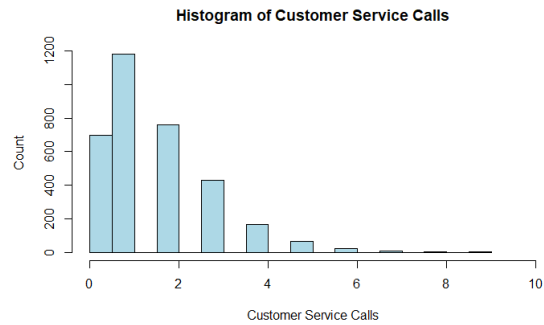
Grafico de Barras Agrupado de Churn e International Plan con referencias

```
barplot(t(counts), col = c("blue", "green"), ylim = c(0, 3300), ylab = "Counts", xlab =
"Churn", main = "International Plan Count by Churn", beside = TRUE)
legend("topright", c(rownames(counts)), col = c("blue", "green"), pch = 15, title =
"Int'l Plan")
box(which = "plot", lty = "solid", col="black")
```



Histograma de las llamadas a atención al cliente (Customer Service Calls)

```
hist(churn$CustServ.Calls, xlim = c(0,10), col = "lightblue", ylab = "Count", xlab =
"Customer Service Calls", main = "Histogram of Customer Service Calls")
```

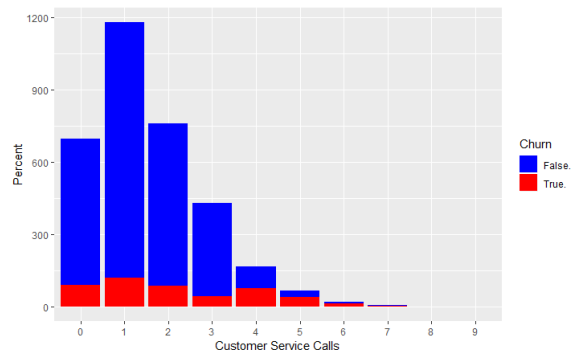


Asegurarse tener instalado el paquete **ggplot2**, sino:


```
install.packages("ggplot2")
# cargarlo
library(ggplot2)
```

Gráficos de Barras Superpuestos

```
ggplot() +
  geom_bar(data = churn, aes(x = factor(churn$CustServ.Calls), fill =
factor(churn$Churn)), position = "stack") +
  scale_x_discrete("Customer Service Calls") +
  scale_y_continuous("Percent") +
  guides(fill=guide_legend(title="Churn")) +
  scale_fill_manual(values=c("blue", "red"))
```



```
ggplot() +
  geom_bar(data=churn, aes(x = factor(churn$CustServ.Calls), fill =
factor(churn$Churn)), position = "fill") +
  scale_x_discrete("Customer Service Calls") +
  scale_y_continuous("Percent") +
  guides(fill=guide_legend(title="Churn"))+
  scale_fill_manual(values=c("blue", "red"))
```



Test T para dos muestras con la variable Int'l Calls

```
# Partitir los datos
churn.false <- subset(churn, churn$Churn == "False.")
churn.true <- subset(churn, churn$Churn == "True.")
# Correr el test t
t.test(churn.false$Intl.Calls, churn.true$Intl.Calls)
```

```
> t.test(churn.false$Intl.Calls, churn.true$Intl.Calls)
```

Welch Two Sample t-test

```
data: churn.false$Intl.Calls and churn.true$Intl.Calls
t = 2.9604, df = 640.64, p-value = 0.003186
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.1243807 0.6144620
sample estimates:
mean of x mean of y
 4.532982  4.163561
```

Grafico de Dispersión de las variables Evening Minutes y Day Minutes, clasificados por Churn

```
plot(churn$Eve.Mins,
     churn$Day.Mins,
     xlim = c(0, 400),
     ylim = c(0, 400),
     xlab = "Evening Minutes",
     ylab = "Day Minutes",
     main = "Scatterplot of Day and Evening Minutes by Churn",
     col = ifelse(churn$Churn == "True.", "red", "blue"))
legend("topright",
     c("True", "False"),
     col = c("red", "blue"),
     pch = 1,
     title = "Churn")
```

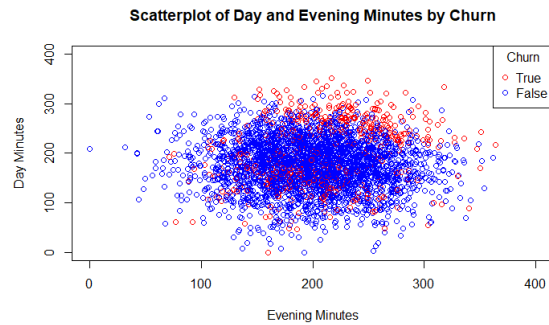




Grafico de dispersión de las variables Day Minutes y Customer Service Calls, coloreado por Churn

```
plot(churn$Day.Mins,
     churn$CustServ.Calls,
     xlim = c(0, 400),
     xlab = "Day Minutes",
     ylab = "Customer Service Calls",
     main = "Scatterplot of Day Minutes and Customer Service Calls by Churn",
     col = ifelse(churn$Churn=="True.", "red", "blue"),
     pch = ifelse(churn$Churn=="True.", 16, 20))
legend("topright",
     c("True", "False"),
     col = c("red", "blue"),
     pch = c(16, 20),
     title = "Churn")
```



Matriz de Dispersion

```
pairs(~churn$Day.Mins+
     churn$Day.Calls+
     churn$Day.Charge)
```

Regresion de Day Charge vs Day Minutes

```
fit <- lm(churn$Day.Charge ~ churn$Day.Mins)
summary(fit)
```

Call:

```
lm(formula = churn$Day.Charge ~ churn$Day.Mins)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.0045935	-0.0025391	0.0004326	0.0024587	0.0045224

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.134e-04	1.711e-04	3.585e+00	0.000341 ***
churn\$Day.Mins	1.700e-01	9.108e-07	1.866e+05	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.002864 on 3331 degrees of freedom

Multiple R-squared: 1, Adjusted R-squared: 1

F-statistic: 3.484e+10 on 1 and 3331 DF, p-value: < 2.2e-16

Correlación con valores p

```
days <- cbind(churn$Day.Mins, churn$Day.Calls, churn$Day.Charge)
MinsCallsTest <- cor.test(churn$Day.Mins, churn$Day.Calls)
MinsChargeTest <- cor.test(churn$Day.Mins, churn$Day.Charge)
CallsChargeTest <- cor.test(churn$Day.Calls, churn$Day.Charge)
round(cor(days), 4)
MinsCallsTest$p.value
MinsChargeTest$p.value
CallsChargeTest$p.value
```

```
> round(cor(days), 4)
      [,1] [,2] [,3]
[1,] 1.0000 0.0068 1.0000
[2,] 0.0068 1.0000 0.0068
[3,] 1.0000 0.0068 1.0000
> MinsCallsTest$p.value
[1] 0.6968515
> MinsChargeTest$p.value
```



```
[1] 0  
> CallsChargeTest$p.value  
[1] 0.6967428
```

Correlación con valores p en forma matricial

```
# Reunir las variables de interes  
corrdata <-  
  cbind(churn$Account.Length,  
        churn$VMail.Message,  
        churn$Day.Mins,  
        churn$Day.Calls,  
        churn$CustServ.Calls)  
# Declarar la matriz  
corrpvalues <- matrix(rep(0, 25), ncol = 5)  
# Llenar la matriz con las correlaciones  
for (i in 1:4) {  
  for (j in (i+1):5) {  
    corrpvalues[i,j] <-  
      corrpvalues[j,i] <-  
        round(cor.test(corrdata[,i], corrdata[,j])$p.value, 4) } }  
round(cor(corrdata), 4)  
corrpvalues
```

```
> round(cor(corrdata), 4)  
      [,1] [,2] [,3] [,4] [,5]  
[1,] 1.0000 -0.0046 0.0062 0.0385 -0.0038  
[2,] -0.0046 1.0000 0.0008 -0.0095 -0.0133  
[3,] 0.0062 0.0008 1.0000 0.0068 -0.0134  
[4,] 0.0385 -0.0095 0.0068 1.0000 -0.0189  
[5,] -0.0038 -0.0133 -0.0134 -0.0189 1.0000  
> corrpvalues  
      [,1] [,2] [,3] [,4] [,5]  
[1,] 0.0000 0.7894 0.7198 0.0264 0.8266  
[2,] 0.7894 0.0000 0.9642 0.5816 0.4440  
[3,] 0.7198 0.9642 0.0000 0.6969 0.4385  
[4,] 0.0264 0.5816 0.6969 0.0000 0.2743  
[5,] 0.8266 0.4440 0.4385 0.2743 0.0000
```

Algunas notas de Árboles de Decisión con R

En el caso del TP_2 se sugiere:

- a) Reducir algunas categorías de las variables, por ejemplo, en *marital_status* y *workclass*
- b) Estandarizar las variables numéricas

Para construir el árbol, es posible emplear el paquete *rpart*

Para dibujar el árbol, es posible utilizar el paquete *rpart.plot*

También es posible construir el árbol utilizando C5.0 (paquete *C50*)

1. Kuhn M, Weston S, Coulter N. 2013. C code for C5.0 by R. Quinlan. C50: C5.0 decision trees and rule-based models. R package version 0.1.0-15. <http://CRAN.R-project.org/package=C50>.
2. Milborrow S. 2012. *rpart.plot*: Plot *rpart* models. An enhanced version of *plot.rpart*. R package version 1.4-3. <http://CRAN.R-project.org/package=rpart.plot>.
3. Therneau T, Atkinson B, Ripley B. 2013. *rpart*: Recursive partitioning. R package version 4.1-3. <http://CRAN.R-project.org/package=rpart>.
4. R Core Team. R: A Language and Environment for Statistical Computing. Vienna, Austria: R Foundation for Statistical Computing; 2012. ISBN: 3-900051-07-0, <http://www.R-project.org/>.